

Chess Analysis Tool

Farees Siddiqui
100-780-513





1. **Defining the Problem**
2. **Chess Vision Parser**
3. **Transformer Chess Chatbot**
4. **Chess Engine**
5. **Chess Engine Fine Tuning**
6. **Results**

For this project I decided to make a Chess analysis tool.

I have created an analysis tool for the game of chess.


The motivation behind this project, lies in its complexity, chess is often known to be a difficult game for humans, and is even more difficult (computationally) for computers. My tool has the following features:

- Computer Vision board parser
 - Used so that the over the board (in person) games can be parsed and analyzed easily without the need for a human to enter the position into a computer manually
- Transformers based Chatbot
 - Used so that the user can ask questions about the current position, as well as ask for assistance, such as the next moves in a certain opening
- Neural Network based Chess Engine
 - Better Generalization than MCTS based approaches
 - Faster Predictions than MCTS based approaches
 - Fine Tuned using self play and Proximal Policy Optimization (PPO)



Chess Vision Parser

- Yolo V5 (You Only Look Once)
- Fine Tuning based on chess pieces dataset
- Data annotation
- Object Detection
- Usage



How Does it Work


- Pre-trained YoloV5 model
- Fine tuned with chess dataset
 - Images of over the board games (top img)
 - Contains annotations of piece labels and positions (bottom img)
- Calculates the FEN string for the position
 - FEN is used to represent a chess position textually
- FEN string is used throughout the tool.



6	0.85009765625	0.782798833819242	0.07177734375	0.13192419825072887
6	0.73291015625	0.08163265306122448	0.05322265625	0.10787172011661808
5	0.70556640625	0.413265306122449	0.0634765625	0.16909620991253643
2	0.77197265625	0.2988338192419825	0.078125	0.1924198250728863
3	0.81640625	0.543002915451895	0.06982421875	0.13119533527696792

Chess Vision Parser Results

- Non Optimal Lighting
- Glare on board and pieces
- Inaccurate detection of square
- Longer Training
- Larger Dataset
- Better Camera Angle
 - Purely overhead would be ideal



Transformer Chess Chatbot

- GPT-2 Model that was fine tuned for chess queries called chessgpt-chat-v1
- For this project, I used the pretrained weights that were available online
- Powered by the transformers architecture
- Training Data:
 - Scraped chess blogs, books, and forumsFiltered and structured into a conversational format
- Answer Questions about chess strategies, openings, and rules

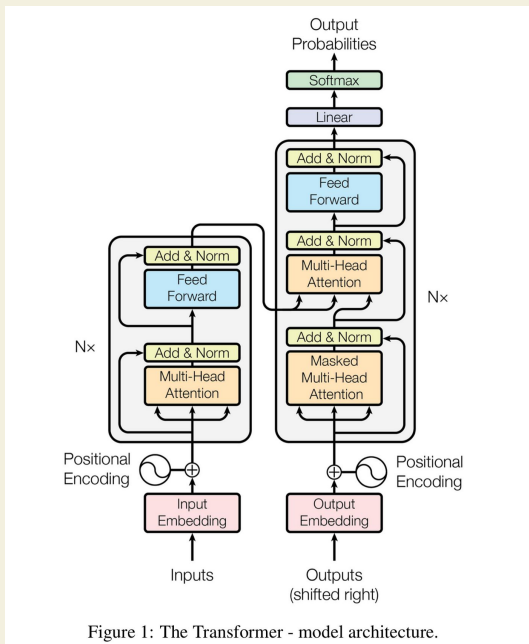


Figure 1: The Transformer - model architecture.

```
prompt = "what is the sicilian defense?"
```

```
Human 1 A: The Sicilian is a defense to 1. e4. It is played by black against 1. e4.
```

```
The Sicilian is not considered a "system" opening per se, but the Kan (4... a6) and Taimanov
```


Dataset & Engine Overview

Dataset Generation

- Extracted 100'000 positions from the Lichess Database
- Data Augmented using stockfish for position evaluations
 - FEN String (Board position)
 - Human Move
 - Stockfish move
 - Stockfish Evaluation of position

Engine Architecture

- Input: Board state encoded as a 12x8x8 tensor
 - 12 chess pieces, 8 rows, 8 columns
- CNN layers to extract spatial features
- Fully Connected Linear layers to predict next moves

```
{
  "fen": "rnbqkbnr/pppppppp/8/8/8/PPPPPPPP/RNBQKBNR w KQkq - 0 1",
  "human_move": "e2e4",
  "stockfish_best_move": "e2e4",
  "stockfish_eval": 40
},
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 8, 8]	6,976
ReLU-2	[-1, 64, 8, 8]	0
Conv2d-3	[-1, 128, 8, 8]	73,856
ReLU-4	[-1, 128, 8, 8]	0
Conv2d-5	[-1, 256, 8, 8]	295,168
ReLU-6	[-1, 256, 8, 8]	0
Flatten-7	[-1, 16384]	0
Linear-8	[-1, 2048]	33,556,480
ReLU-9	[-1, 2048]	0
Linear-10	[-1, 1024]	2,098,176
ReLU-11	[-1, 1024]	0
Linear-12	[-1, 64]	65,600
Linear-13	[-1, 512]	8,421,888
ReLU-14	[-1, 512]	0
Linear-15	[-1, 64]	32,832

Total params: 44,550,976
 Trainable params: 44,550,976
 Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 0.62
 Params size (MB): 169.95
 Estimated Total Size (MB): 170.57

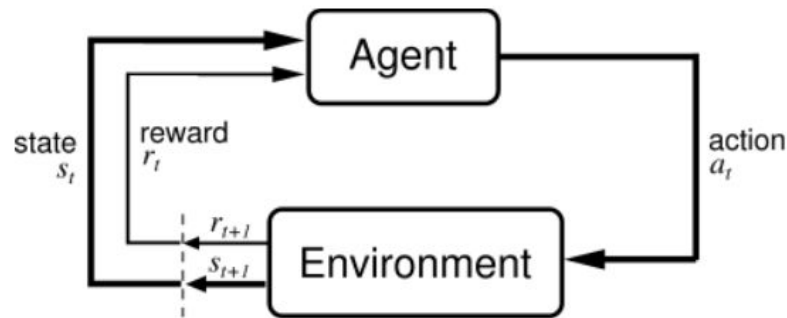
Reinforcement Learning & PPO?

RL Overview

- **Agent:** Learns a policy to decide actions based on the current state
- **Environment:** Provides feedback by returning a new state and a reward after the agent performs an action
- **State Space:** represents the state space of the environment
- **Action Space:** represents all possible actions that an agent can perform in a given state

PPO Overview

- Policy Optimization method that helps to stabilize the learned policy
- Optimizes a policy using small updates
 - Avoids drastic changes
- Clipped loss function is used to prevent the new policy from diverging too far from the old one



Algorithm 1 PPO-Clip

- 1: Input: initial policy parameters θ_0 , initial value function parameters ϕ_0
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_k = \pi(\theta_k)$ in the environment.
- 4: Compute rewards-to-go \hat{R}_t .
- 5: Compute advantage estimates, \hat{A}_t (using any method of advantage estimation) based on the current value function V_{ϕ_k} .
- 6: Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left(\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} A^{\pi_{\theta_k}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

- 7: Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 8: **end for**

RL Fine Tuning for NN Engine

Objective

- Use Proximal Policy Optimization (PPO) to train a reinforcement learning agent to play chess
- Improve the neural network based model by having it play against stockfish and learn from it

Custom Chess Environment

- Encode the chessboard as a 12x8x8 tensor
 - To conform to NN input
- 64x64 action space for all possible moves
 - Possible moves != valid moves
 - Possible moves are any combination of moving from a square to another
 - Use stockfish as an opponent to learn from

